


Large Scale Training

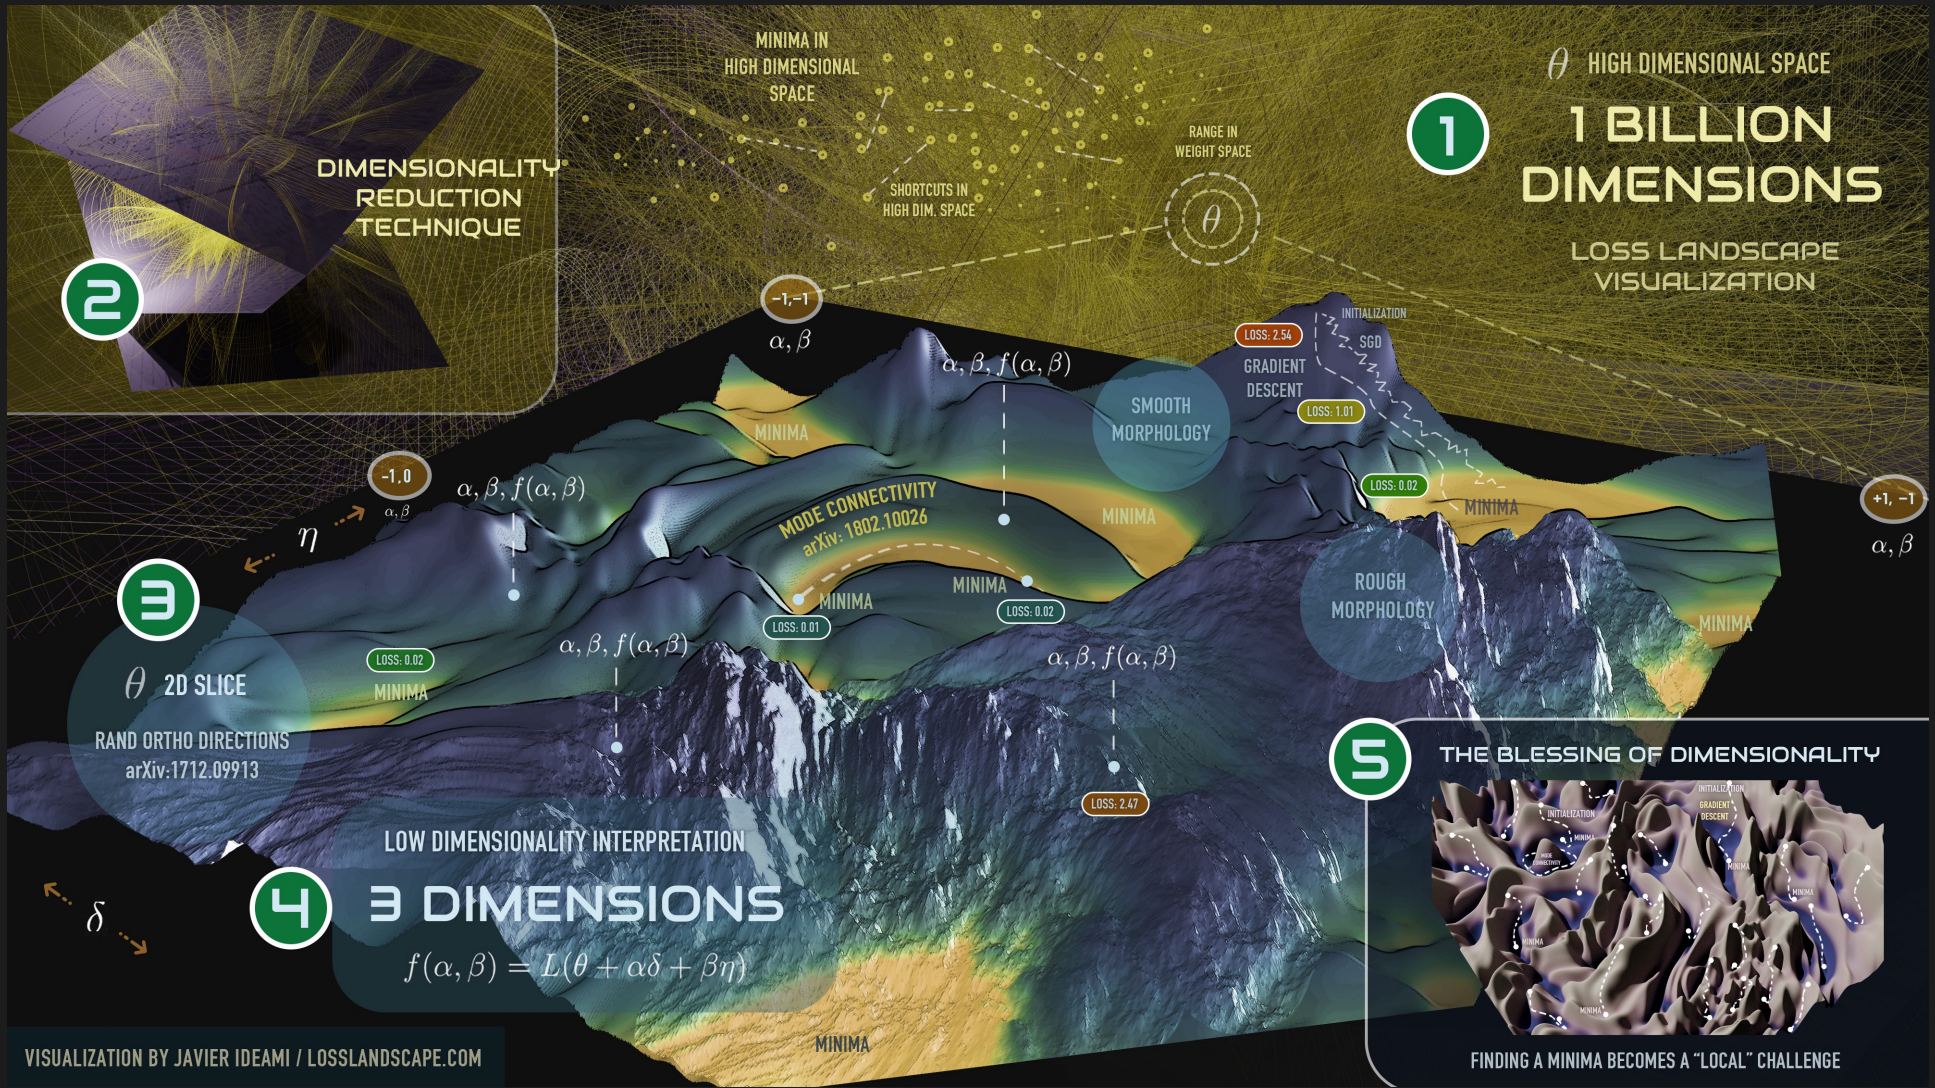
 [argonne-lcf / ai-science-training-series](#)

 Sam Foreman

2022-11-01

Why Distributed Training?

- Large batches may not fit in GPU memory
- Splitting data across workers → larger batch size
- Smooth loss landscape
- Improved gradient estimators
- Less iterations needed for same number of epochs
 - May need to train for more epochs if another change is not made
 - e.g. scaling learning rate
- See [Large Batch Training of Convolutional Networks](#)



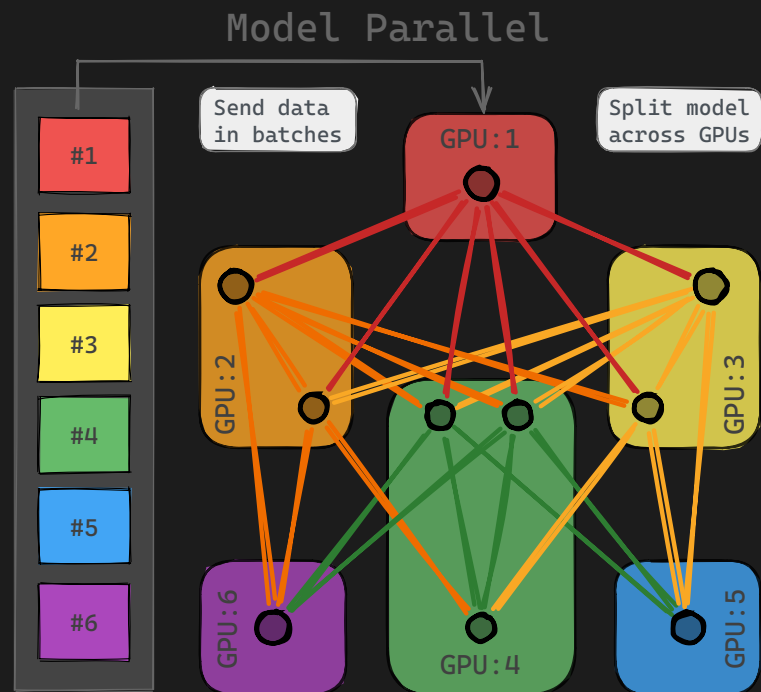
Recent Progress

Year	Author	Batch Size	Processor	DL Library	Time	Accuracy
2016	He et al. [1]	256	Tesla P100 x8	Caffe	29 Hrs	75.3%
	Goyal et al. [2]	8192	Tesla P100	Caffe 2	1 hour	76.3%
	Smith et al. [3]	8192 → 16,384	full TPU pod	TensorFlow	30 mins	76.1%
	Akiba et al. [4]	32,768	Tesla P100 x1024	Chainer	15 mins	74.9%
	Jia et al. [5]	65,536	Tesla P40 x2048	TensorFlow	6.6 mins	75.8%
	Ying et al. [6]	65,536	TPU v3 x1024	TensorFlow	1.8 mins	75.2%
	Mikami et al. [7]	55,296	Tesla V100 x3456	NNL	2.0 mins	75.29%
2019	Yamazaki et al	81,920	Tesla V100 x 2048	MXNet	1.2 mins	75.08%

Model Parallel Training

Model Parallel Training


- Split up network over multiple workers
 - Each receives disjoint subset
 - All communication associated with subsets are distributed
- Communication whenever dataflow between two subsets
- Typically **more complicated** to implement than data parallel training
- Suitable when the model is too large to fit onto a single device (CPU / GPU)



Model Parallel Training

$$y = w_0 * x_0 + w_1 * x_1 + w_2 * x_2$$

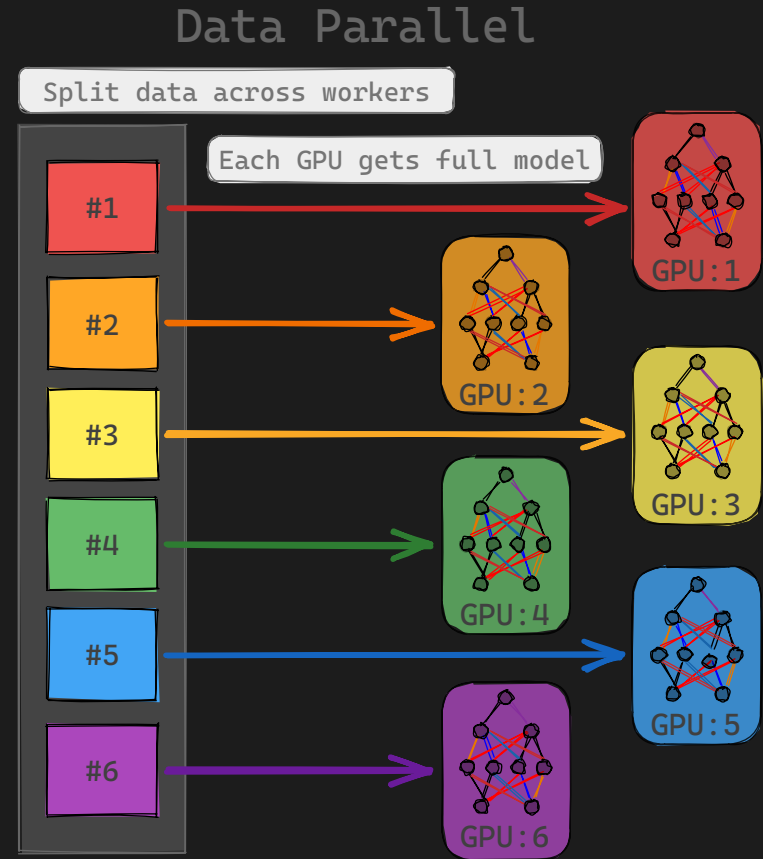
1. Compute $y_0 = w_0 * x_0$ and send to \rightarrow GPU1
2. Compute $y_1 = y_0 + w_1 * x_1$ and send to \rightarrow GPU2
3. Compute $y = y_1 * w_2 * x_2$ ✓



Data Parallel Training

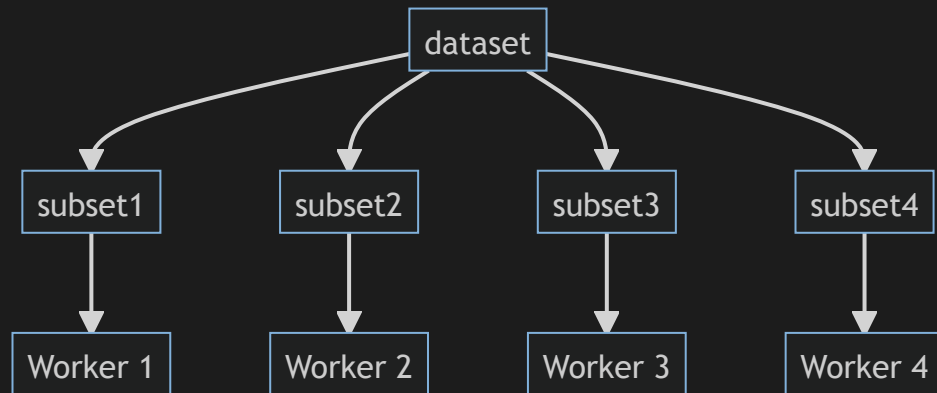
Data Parallel Training

- Each worker has identical copy of complete model
- Each Worker computes the corresponding loss and gradients w.r.t **local** data
- Before updating parameters, loss and gradients averaged across workers
- Typically easier / simpler to implement



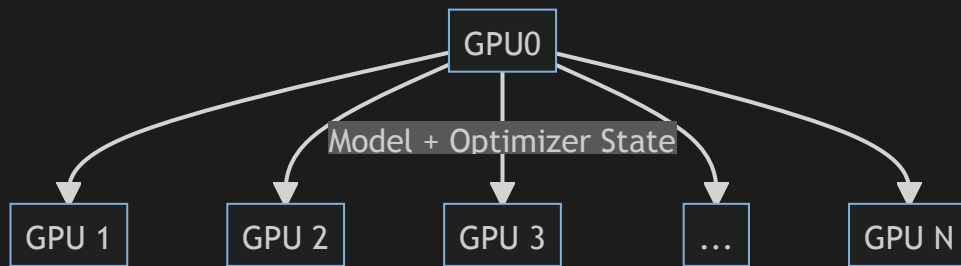
Data Parallel Training

- Each worker has identical copy of model
- **Global batch of data split across workers**
- Loss + Grads averaged across workers before updating parameters



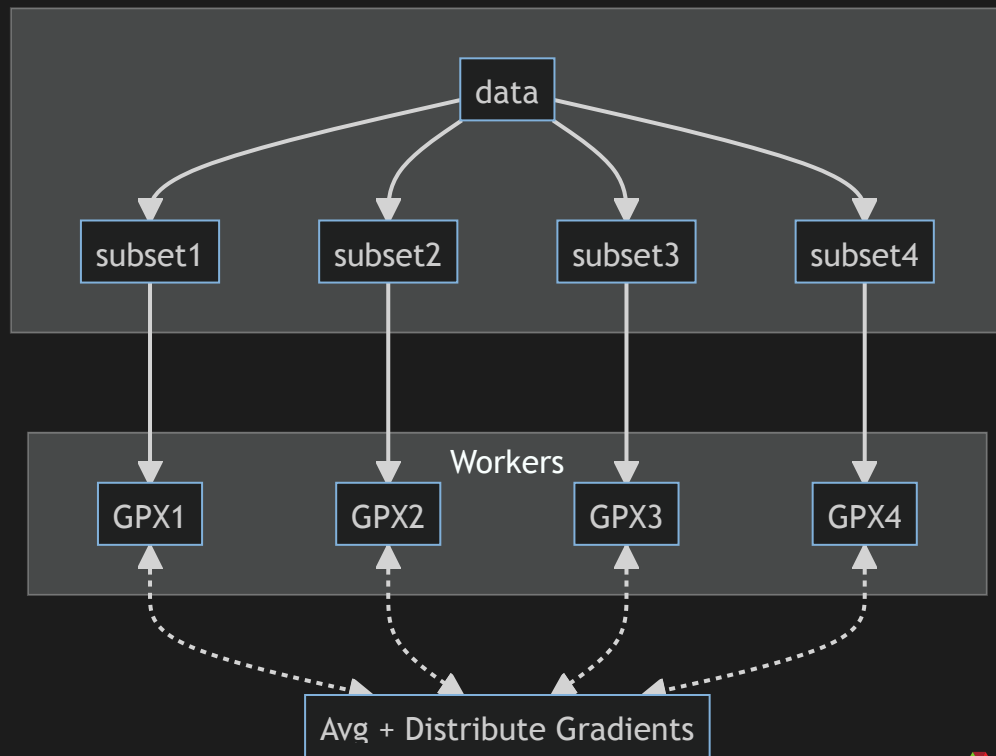
Broadcast Initial State

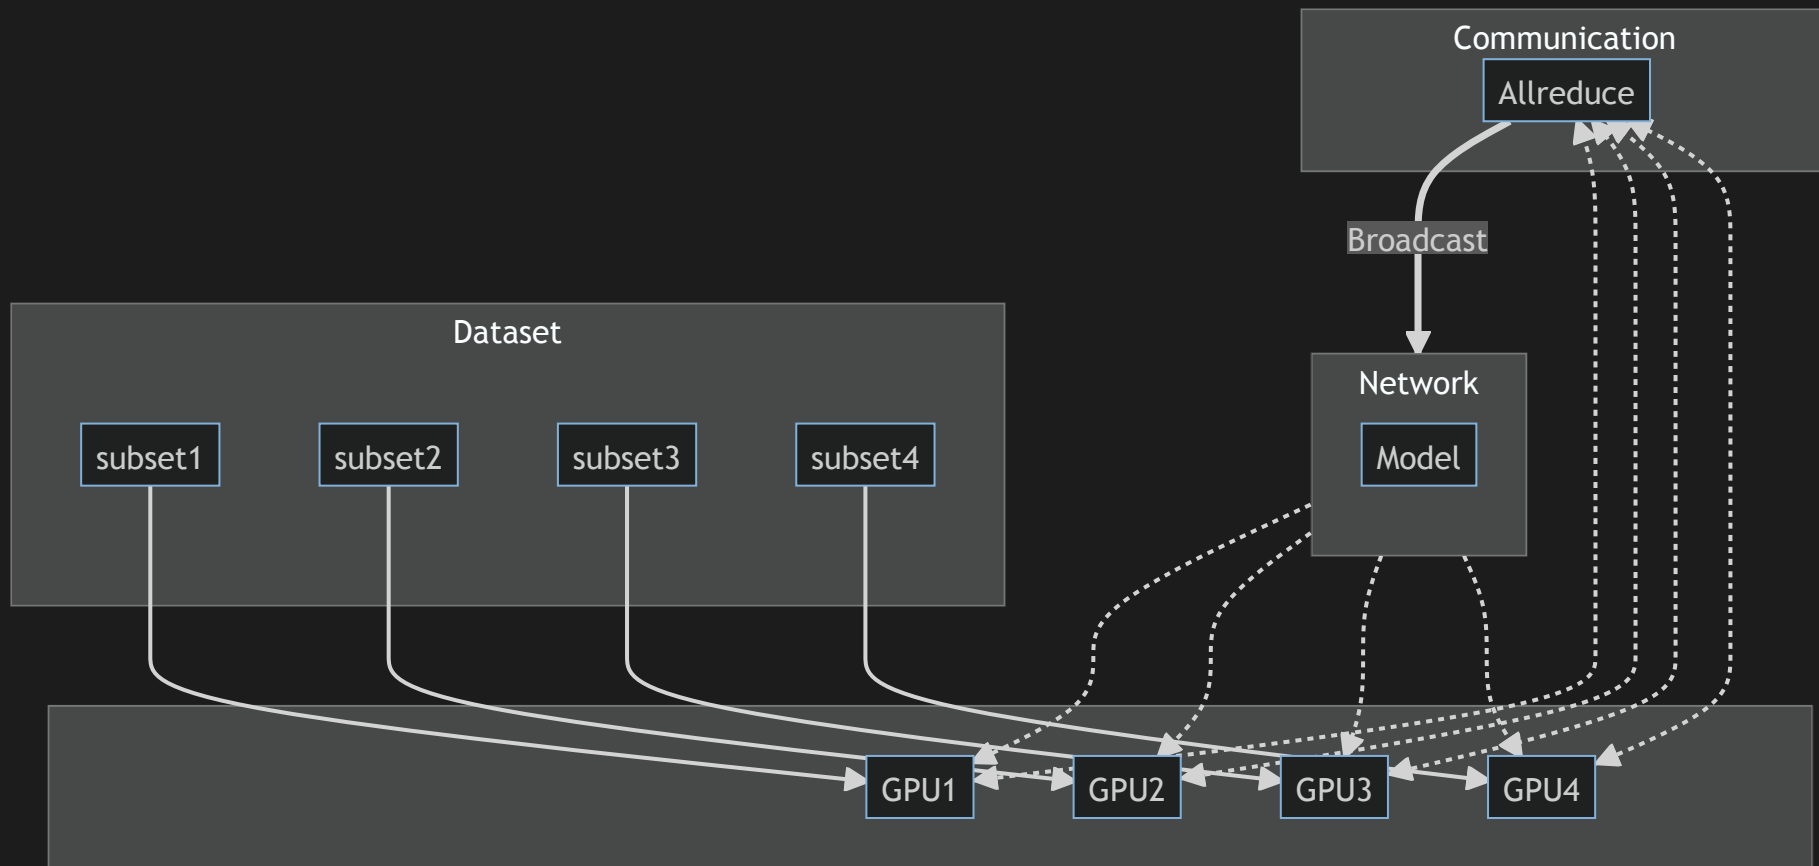
- At the start of training (or when loading from a checkpoint), we want all of our workers to be initialized consistently
 - Broadcast** the model and optimizer states from `hvd.rank() == 0` worker



Data Parallel Training

- Disjoint subsets of a neural network are assigned to different devices
- Each worker receives:
 - **identical copy of model**
 - **unique subset of data**





TensorFlow + Horovod

- Set one GPU per process ID (`hvd.local_rank()`)

```
gpus = tf.config.experimental.list_physical_devices('GPU')
for gpu in gpus:
    tf.config.experimental.set_memory_growth(gpu, True)
if gpus:
    local_rank = hvd.local_rank()
    tf.config.experimental.set_visible_devices(gpus[local_rank], 'GPU')
```

Scale the Learning Rate

- Scale the learning rate by the number of workers to account for the increased batch size

```
import horovod.tensorflow as hvd
optimizer = tf.optimizers.Adam(lr_init * hvd.size())
```


Thank you!

- Organizers
- ALCF Data Science & Operations
- Feel free to reach out!



Acknowledgements

This research used resources of the Argonne Leadership Computing Facility, which is a DOE Office of Science User Facility supported under Contract DE-AC02-06CH11357.

